

# CS-344 ASSIGNMENT-3

## GROUP-09

Patnana Sai Ashrrith - 190101061  
Amancha Jagruth - 190101010  
Bhukya Bharath - 190101024  
Adduri Sri Sai Datta - 190101003

## PART-A :- LAZY MEMORY ALLOCATION

### Task -1 : eliminate allocation from sbrk()

The extra memory for a current process is obtained by using sbrk() system call, but the growproc() line in the sbrk system call is commented out here. The new sbrk(n) will increment the memory size parameter of the process by n (**line 54**) and return the previous size without actually increasing the memory. When this process tries to access the extra memory, a page fault occurs. Thus generating the T\_PGFLT trap to the kernel.

```
44
45 int
46 sys_sbrk(void)
47 {
48     int addr;
49     int n;
50
51     if(argint(0, &n) < 0)
52         return -1;
53     addr = myproc()->sz;
54     myproc()->sz+=n;
55     //if(growproc(n) < 0)
56         // return -1;
57     return addr;
58 }
59
```

### Task -2 : Lazy allocation

We need to restrain giving memory as soon as it is requested. Rather, we give the memory when it tries to access. This is Lazy Memory Allocation. The page fault is handled by **PGFLT\_handler()** in trap.c.

Working of the PGFLT\_handler function:

- This function is called when a T\_PGFLT trap is generated.
- Now, rcr2() returns the virtual address at which the page fault occurs.
- rounded\_addr points to the starting address to the page where this virtual address resides.( This rounded address is generated using PGROUNDDOWN macro.
- Then we call kalloc() which returns a free page from a linked list of free pages (freelist inside kmem) in the system.
- We have a physical page at our disposal. Now we need to map it to the virtual address rounded\_addr which is done using mappages(), if no free pages are available, kalloc will return 0, then we will exit the function returning -1.
- To use mappages() in trap.c, we removed the static keyword in front of it in vm.c and declared its prototype in trap.c.
- mappages() takes the page directory (myproc()->pgdir) of the current process, virtual address of the start of the data where the page fault occurs, size of the data, physical memory at which the physical page resides (we give this parameter by using V2P macro which converts our virtual address to physical address by subtracting KERNBASE from it) and permissions corresponding to the page table entry as parameters.

```
96         break;
97     case T_PGFLT:
98         if(PGFLT_handler()<0){
99             cprintf("Could not allocate page. Sorry.\n");
100             panic("trap");
101         }
102     break;
103
```

```
17 int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm);
18
19 int PGFLT_handler(){
20     int addr=rcr2();
21     int rounded_addr = PGROUNDDOWN(addr);
22     char *mem=kalloc();
23     // if there is memory available, we will allocate it to the process
24     if(mem!=0){
25         memset(mem, 0, PGSIZE);
26         // maps the physical address to the virtual address
27         if(mappages(myproc()->pgdir, (char*)rounded_addr, PGSIZE, V2P(mem), PTE_W|PTE_U)<0)
28             return -1;
29         return 0;
30     } else
31         return -1;
32 }
33
```

- In `mappages()` loop runs until all the pages from the first to last have been loaded successfully. For every page, it loads it into the page table using `walkpgdir()`.

Test cases :

Everything is working fine as the page fault is taken care of. Basic commands are running as shown in the below-left image..

```
60 int // removed static
61 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
62 {
63     char *a, *last;
64     pte_t *pte;
65
66     a = (char*)PGROUNDDOWN((uint)va);
67     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
68     for(;;){
69         if((pte = walkpgdir(pgdir, a, 1)) == 0)
70             return -1;
71         if(*pte & PTE_P)
72             panic("remap");
73         *pte = p char *last TE_P;
74         if(a == last)
75             break;
76         a += PGSIZE;
77         pa += PGSIZE;
78     }
79     return 0;
80 }
```

## PART-B :- XV6 Memory

**Q1: How does the kernel know which physical pages are used and unused?**

- XV6 keeps a linked list of free pages called `kmem` in `kalloc.c`.
- The address to the next free page is stored in the first address of the present free page. This is one of the advantages of maintaining a linked list of free pages.
- These lists were empty initially, `xv6` calls `kinit` from `Main` which then adds 4MB of free pages to the list

**Q2: What data structures are used to answer this question?**

- Linked lists are used

**Q3: Where do these reside?**

- These linked lists reside in `kalloc.c`
- Every node of these linked lists is a structure defined in `kalloc.c` (`struct run`)

**Q4: Does `xv6` memory mechanism limit the number of user processes?**

- The number of user processes are limited in `xv6`, due to the limit in size of `ptable`.

**Q5: If so, what is the lowest number of processes `xv6` can 'have' at the same time (assuming the kernel requires no memory whatsoever)?**

- There is one process named `initproc` when the `xv6` system boots up.
- Process can have a virtual address space of `2GB - KERNBASE` and maximum physical memory of 240MB
- `PHYSTOP`
- 1 process can take up all 240MB of physical memory, the lowest number of processes in `xv6` is 1.
- There can't be zero processes after boot, since all user interactions need to be done using user processes which are forked from `initproc`.

```
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat        2 3 16280
echo       2 4 15136
forktest  2 5 9440
grep       2 6 18500
init       2 7 15720
kill       2 8 15164
ln         2 9 15020
ls         2 10 17648
mkdir      2 11 15264
rm         2 12 15240
sh         2 13 27876
stressfs   2 14 16156
usertests  2 15 67260
wc         2 16 17016
zombie     2 17 14832
console    3 18 0
$ echo ho
ho
$
```

## Task-1: Kernel processes

In `proc.c`, `create_kernel_process()` function is created. During the whole time, the kernel process will remain in kernel mode. So, there is no need to initialize its trapframe, user space and the user section of the page table. Address of the instruction is stored in the `eip` register. We set the `eip` value of the context to the entry point as we want the process to start executing at the entry point.

`Allocproc` assigns the process a spot in the table.

`SetupKvm` sets the kernel part of the process table which maps virtual addresses above `KERNBASE` to physical addresses between 0 and `PHYSTOP`.

## Task-2: swapping out mechanism

A circular queue is created as `swap_req`.

`Swap_out_que` is a specific queue, it holds the processes with swap out requests.

Functions corresponding to `swap_req` are also created as `swap_req_push()` and `swap_req_pop()`.

Queue can be accessed with a lock that we have initialised in `pinit`.

We added prototypes in `def.h` as we need the queue and functions relating to it in other files too.

Now, in `allocuvm()` which is called by `growproc()`, whenever `kalloc()` is not able to allocate free pages. Then we create a `SWAP_OUT_PROCESS` using `create_kernel_process`.

```
mem = kalloc();
if(mem == 0){
    // cprintf("allocuvm out of memory\n");
    deallocuvm(pgdir, newsz, oldsz);

    //SLEEP
    myproc->state = SLEEPING;
    acquire(&swapsleeplock);
    myproc->chan = swapsleep;
    swapsleepcount++;
    release(&swapsleeplock);

    swap_req_push(myproc(), &swap_out_req);
    if(!SOP_PRESENT){
        // if condition to make sure that only one SWAP_OUT_PROCESS exists at a given time.
        SOP_PRESENT = 1;
        create_kernel_process("SWAP_OUT_PROCESS", &SWAP_OUT_PROCESS);
    }

    return 0;
}
```

```
450
451 void create_kernel_process(const char *name, void (*entrypoint)()){
452
453     struct proc *p = allocproc();
454
455     if(p == 0)
456         panic("create_kernel_process failed");
457
458     //Setting up kernel page table using setupkvm
459     if((p->pgdir = setupkvm()) == 0)
460         panic("setupkvm failed");
461
462     //This is a kernel process. Trap frame stores user space registers. We don't need to initialise tf.
463     //Also, since this doesn't need to have a userspace, we don't need to assign a size to this process.
464
465     //eip stores address of next instruction to be executed
466     p->context->eip = (uint)entrypoint;
467
468     safestrcpy(p->name, name, sizeof(p->name));
469
470     acquire(&ptable.lock);
471     p->state = RUNNABLE;
472     release(&ptable.lock);
473
474 }
```

```
170
171 struct swap_req{
172     struct spinlock lock; // lock to restrict access of this swap request queue
173     struct proc* queue[NPROC];
174     int start;
175     int end;
176 };
177
178 // request queue for swapping out requests
179 struct swap_req swap_out_req;
180 // request queue for swapping in requests
181 struct swap_req swap_in_req;
182
183 struct proc* swap_req_pop(struct swap_req *q){
184
185     acquire(&q->lock);
186     if(q->start == q->end){
187         release(&q->lock);
188         return 0;
189     }
190     struct proc *p = q->queue[q->start];
191     (q->start)++;
192     (q->start) %= NPROC;
193     release(&q->lock);
194
195     return p;
196 }
197
198 int swap_req_push(struct proc *p, struct swap_req *q){
199
200     acquire(&q->lock);
201     if((q->end+1)%NPROC == q->start){
202         release(&q->lock);
203         return 0;
204     }
205     q->queue[q->end] = p;
206     q->end++;
207     (q->end) %= NPROC;
208     release(&q->lock);
209
210     return 1;
211 }
```

Now, we need to change the process state to sleeping, and its channel is set to `swapsleep`.

So, a current process is added to swap out the request queue. (`swap_out_req`).

Now here, `SOP_PRESENT` ensures that only one swap out process exists at a given moment. This bit is set to 0 in `SWAP_OUT_PROCESS` function,

```

61 kfree(char *v)
62 {
63     struct run *r;
64
65     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
66         panic("kfree");
67
68     // Fill with junk to catch dangling refs.
69     // memset(v, 1, PGSIZE);
70     for(int i=0; i<PGSIZE; i++) v[i] = 1;
71
72     if(kmem.use_lock)
73         acquire(&kmem.lock);
74     r = (struct run*)v;
75     r->next = kmem.freelist;
76     kmem.freelist = r;
77     if(kmem.use_lock)
78         release(&kmem.lock);
79
80     //Wake up processes sleeping on swapsleep channel.
81     if(kmem.use_lock)
82         acquire(&swapsleeplock);
83     if(swapsleepcount) {
84         wakeup(swapsleep);
85         swapsleepcount--;
86     }
87     if(kmem.use_lock)
88         release(&swapsleeplock);
89 }
90

```

Now, if memory is allocated it simply maps the new pages to the virtual addresses using mappages. (not included in the screenshot).

Kalloc.c:-

We created a mechanism which wakes up all the sleeping processes in swapsleep whenever free pages are available.

Kfree in kalloc.c is edited in this way:

Processes that were preempted because of lack of availability of pages were sent to sleeping on the swapsleep.

wakeup() system call wakes up all the processes currently sleeping in swapsleep. Here we also commented on the memset line(69) as we do not want to risk erasing data before we move it to the hard disk.

Now, lets see the **SWAP\_OUT\_PROCESS** function.

Process runs while the loop till the swap out requests queue is not empty.

Loops runs with popping the first process in the queue and uses the LRU policy to find a victim page in the page table.

We iterate through each entry in the process table(pgdir) which thereby extracts the physical address for the secondary page.

And for the secondary page table, we iterate among the page table and look for accessed bit(A) on each of the entries.

When the secondary page table entry is found with an accessed bit unset, it chooses this entry's physical page number as the victim page.

This page then swapped out and stored to drive.

Pid and virtual address of the page to be eliminated to name the file storing this page.

New function int2str copies integers into a given string.

When no requests are left, kernel process is suspended:

When the queue is empty, loop breaks and process suspension is initiated. We couldn't clear their kstack from inside the process while exiting the kernel processes as they won't know which process to execute next.

```

213 void SWAP_OUT_PROCESS() {
214
215     acquire(&swap_out_req.lock);
216     while(swap_out_req.start != swap_out_req.end){
217         struct proc *p = swap_req_pop(&swap_out_req);
218
219         pde_t *pgdir = p->pgdir;
220         for(int i=0; i<NPENTRIES; i++){ // going through the page directory entries.
221
222             //skip page table if accessed. chances are high, not every page table was accessed.
223             if(pgdir[i] & PTE_A) continue;
224
225             pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(pgdir[i]));
226             for(int j=0; j<NPENTRIES; j++){ // going through the the page table entries
227
228                 //Skip if found
229                 if((pgtab[j]&PTE_A) || !(pgtab[j]&PTE_P)) continue;
230
231                 pte_t *pte = (pte_t*)P2V(PTE_ADDR(pgtab[j]));
232
233                 //for file name
234                 int pid = p->pid;
235                 // file name contains virtual address of the swapping out which helps SWAP_IN_PROCESS
236                 // to swap in a particular page fault at a given address by the process.
237                 int virt_addr = ((1<<22)*i)+((1<<12)*j);
238                 //file name
239                 char c[50];
240                 num_to_str(pid,c);
241                 int x = strlen(c);
242                 c[x] = '-';
243                 num_to_str(virt_addr,c+x+1);
244                 memset(&pgtab[j], 0, sizeof(pgtab[j]));
245
246                 //mark this page as being swapped out.
247                 pgtab[j] = ((pgtab[j])^(0x000));
248
249                 break;
250             }
251         }
252     }
253
254     release(&swap_out_req.lock);
255
256     struct proc *p;
257     if((p=myproc()) == 0)
258         panic("swap out process");
259
260     SOP_PRESENT = 0; // setting it zero so that another new SWAP_OUT_PROCESS can be created.
261     p->parent = 0;
262     p->name[0] = '\0';
263     p->killed = 0;
264     p->state = UNUSED; // Killing this swapping out process.
265     sched(); // calling scheduler.
266 }
267

```



So, we need to clear the kstack from outside the process. We first preempt the process process and wait for the scheduler to find this process. If a kernel process in UNUSED states is found by scheduler, it clears process kstack and name. This is identified by checking its name in which the first character was changed to “\*” when the process ended. And when scheduler selects a process, its access bit of every PDE and PTE are reset. We defined another bit flag named PTE\_A of value 0x020 which marks an entry accessed or not.

```

706 //If the swap processes have stopped running, free its stack and name.
707 if(p->state==UNUSED && p->name[0]=='*'){
708     kfree(p->kstack);
709     p->kstack = 0;
710     p->name[0] = 0;
711     p->pid = 0;
712 }
713
714 if(p->state != RUNNABLE)
715     continue;
716
717 // we will reset the access bit of the selected process as we will just mark a recently used if it used in the last quantum of the process.
718 for(int i=0; i<NPDETRIES; i++){
719     //If PDE was accessed
720
721     if(((p->pgdir)[i]&PTE_P && ((p->pgdir)[i]&PTE_A){
722         pte_t* pgtab = (pte_t*)P2V(PTE_ADDR((p->pgdir)[i]));
723         for(int j=0; j<NPTETRIES; j++){
724             if(pgtab[j] & PTE_A){
725                 pgtab[j] ^= PTE_A;
726             }
727         }
728         ((p->pgdir)[i]) ^= PTE_A;
729     }
730 }

```

### Task-3 : Swapping in Mechanism

- When the kernel detects a page fault, it must check if the cause of this page fault is in the swapping out mechanism.

In Task 2, if we swapped out a page we set its page table entries a bit of 7th order(0x080 also PTE\_PS). So to check if the page was swapped out we check its 7th bit if it is set we call swap\_in\_process else exit.

- Swapping in the function runs a while loop until Swap\_in\_queue is not empty. In the loop, it pops a process from the queue and extracts its pid and addr value to get the file name. Then, it creates the filename in a string called “c” using num\_to\_string(). Then, it uses file\_open() to open this file in read only mode (O\_RDONLY) with file descriptor fd. We then

```

19 void PGFLT_handler() {
20     int addr=rcr2();
21     struct proc *p = myproc();
22     acquire(&swpinlock);
23     sleep(p, &swpinlock);
24     pde_t *pde = &(p->pgdir)[PDX(addr)];
25     pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
26
27     if((pgtab[PTX(addr)]&0x080){
28         //This means that the page was swapped out.
29         //virtual address for page
30         // storing the address where page fault occurs. This is later used to swap in the respective file .swp file
31         p->PGFLT_addr = addr;
32         swap_req_push(p,&swp_in_req);
33         if(!SIP_PRESENT){
34             SIP_PRESENT = 1;
35             create_kernel_process("SWAP_IN_PROCESS", &SWAP_IN_PROCESS);
36         }
37     }
38     else exit();
39 }
40

```

```

294 void SWAP_IN_PROCESS() {
295
296     acquire(&swp_in_req.lock);
297     while(swap_in_req.start != swap_in_req.end){
298         struct proc *p = swap_req_pop(&swp_in_req);
299
300         int pid = p->pid;
301         int virt_addr = PTE_ADDR(p->PGFLT_addr);
302
303         char c[50];
304         num_to_str(pid,c);
305         int x = strlen(c);
306         c[x] = '-';
307         num_to_str(virt_addr,c+x+1); // getting the page which existed at this va before getting swapped out.
308         safestrcpy(c+strlen(c),".swp",5);
309
310         int fd = open_file(c,O_RDONLY);
311         if(fd<0){
312             release(&swp_in_req.lock);
313             printf("could not find page file in memory: %start\n", c);
314             panic("SWAP_IN_PROCESS");
315         }
316         char *mem = kalloc();
317         read_file(fd,PGSIZE,mem);
318
319         if(mappages(p->pgdir, (void *)virt_addr, PGSIZE, V2P(mem), PTE_W|PTE_U)<0){
320             release(&swp_in_req.lock);
321             panic("mappages");
322         }
323         wakeup(p);
324     }
325 }

```

allocate a free frame (mem) to this process using kalloc. We read from the file with the fd file descriptor into this free frame using read2. We then make mappages available to proc.c by removing the static keyword from it in vm.c and then declaring a prototype in proc.c. We then use mappages to map the page

```

325
326     release(&swap_in_req.lock);
327     struct proc *p;
328     if((p=myproc()) == 0)
329         panic("SWAP_IN_PROCESS");
330
331     SIP_PRESENT = 0; // resetting the value so a new SWAP_OUT_PROCESS may be created.
332     p->parent = 0;
333     p->name[0] = '*';
334     p->killed = 0;
335     p->state = UNUSED;
336     sched(); // calling the scheduler.
337 }
338

```

corresponding to addr with the physical page that got using kalloc and read into (mem). Then we wake up, the

process for which we allocated a new page to fix the page fault using wakeup. Once the loop is completed.

Suspending kernel process when no requests are left: When the queue is empty, the loop breaks and suspension of the process is initiated just like in task 2.

## Task-4 : Sanity Test

### Observation in implementation:

```

1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  int allocnum(int n){
6      return n*n - 4*n + 1;
7  }
8
9  int
10 main(int argc, char* argv[]) {
11
12     for(int i=0; i<20; i++){
13         if(fork() == 0){
14             printf(1, "Child %d\n", i+1);
15             printf(1, "   S.no   Matched   Error\n");
16             printf(1, "-----\n\n");
17
18             for(int j=0; j<10; j++){
19                 int *a = malloc(4096);
20                 for(int k=0; k<1024; k++) a[k] = allocnum(k);
21
22                 int Matched_B = 0;
23                 for(int k=0; k<1024; k++) if(a[k] == allocnum(k)) Matched_B += 4;
24
25                 if(j<9) printf(1, "   %d   %dB   %dB\n", j+1, Matched_B, 4096 - Matched_B);
26                 else printf(1, "   %d   %dB   %dB\n", j+1, Matched_B, 4096 - Matched_B);
27             }
28             printf(1, "\n");
29
30             exit();
31         }
32     }
33 }
34
35 while(wait()!=-1);
36 exit();
37
38
39

```

- The main process will fork 20 child processes.
- Each child process executes a loop with 10 iterations.
- At each iteration the process will allocate 4KB of memory using malloc system call.
- next it will fill the memory with values obtained from allocnum function (which returns  $n^2-4n-1$ ).

- Num\_of\_bytes\_matched variable counts number of bytes that are matched in the stored value and the calculated value.

Testing using different values of the PHYSTOP values

**Case 1 :** PHYSTOP value (0xE000000-224MB)

\$ sanity		
Child 1		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B
Child 2		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B
Child 3		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B

**Case 2 : PHYSTOP value (0x0800000)**

we get the same result as case1 even after reducing the memory size.

```
$ sanity
```

Child 1		
S.no	Matched	Error
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B

  

Child 2		
S.no	Matched	Error
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B

  

Child 3		
S.no	Matched	Error
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B



**Case 3 : PHYSTOP value (0x0400000-4MB)**

\$ sanity		
Child 1		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B
Child 2		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B
Child 3		
S.no	Matched	Error
-----		
1	4096B	0B
2	4096B	0B
3	4096B	0B
4	4096B	0B
5	4096B	0B
6	4096B	0B
7	4096B	0B
8	4096B	0B
9	4096B	0B
10	4096B	0B
Child 4		
S.no	Matched	Error

Here we use 4MB because of the minimum memory needed by memory to execute Kinit1. The obtained output is the same as the previous output representing the implementation is correct.